

lab10-miniproject

Meha Thakur PID 16020450

Table of contents

1. Importing candy data	1
2. What is your favorite candy?	2
4. Taking a look at pricepercent	10
5 Exploring the correlation structure	12
6. Principal Component Analysis	13

Happy Halloween!!

1. Importing candy data

```
#reading in data and loading packages
candy <- read.csv("candy-data.csv", row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat	crisped	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0
	hard	bar	pluribus	sugar	percent	pricepercent	winpercent		
100 Grand	0	1	0	0.732		0.860	66.97173		
3 Musketeers	0	1	0	0.604		0.511	67.60294		
One dime	0	0	0	0.011		0.116	32.26109		
One quarter	0	0	0	0.011		0.511	46.11650		
Air Heads	0	0	0	0.906		0.511	52.34146		
Almond Joy	0	1	0	0.465		0.767	50.34755		

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(skimr)
library(rlang)
library(ggplot2)
```

The functions dim(), nrow(), table() and sum() may be useful for answering the first 2 questions.

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

[1] 85

There are 85 candy types

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity==1)
```

[1] 38

There are 38 fruity candy types in the dataset

2. What is your favorite candy?

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

```
#example -- candy["Twix", ]$winpercent  
candy["Reese's Peanut Butter cup", ]$winpercent
```

```
[1] 84.18029
```

I love Reese's peanut butter cups, winpercent value 84.18029

Q4. What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat", ]$winpercent
```

```
[1] 76.7686
```

76.7686

Q5. What is the winpercent value for "Tootsie Roll Snack Bars"?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

49.6535

Hin (6 and 7): look at the "Variable type" print out from the skim() function. Most variables (i.e. columns) are on the zero to one scale but not all. Some columns such as chocolate are exclusively either zero or one values.

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

```
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12

Column type frequency:

numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmond	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

winpercent appears to be the only row with a mean more than 1, or without 0-1 values.

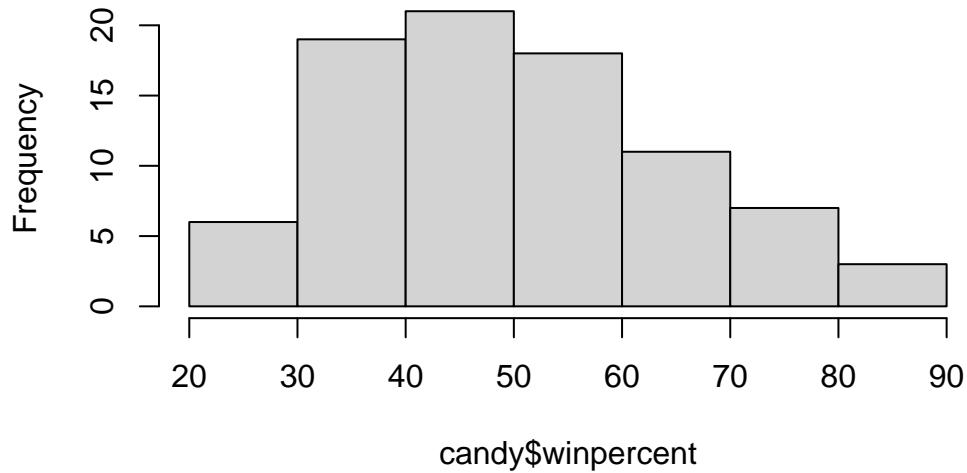
Q7. What do you think a zero and one represent for the candy\$chocolate column?

0 and 1 likely represent yes or no. For example, is this candy specified (for example, Air Heads) a chocolate. It is not, according to the dataset given by its value of zero. I agree with this evaluation.

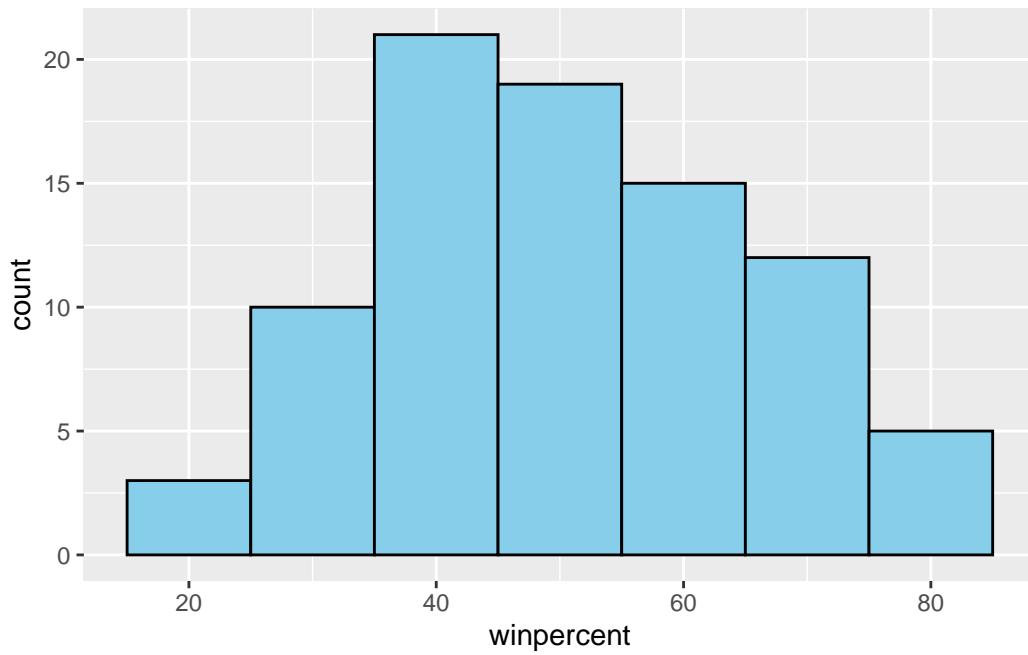
Q8. Plot a histogram of winpercent values

```
hist(candy$winpercent)
```

Histogram of candy\$winpercent



```
ggplot(candy,aes(winpercent))+  
  geom_histogram(binwidth=10,fill="skyblue", col="black")
```



Q9. Is the distribution of winpercent values symmetrical?

```
mean(candy$winpercent)
```

```
[1] 50.31676
```

```
median(candy$winpercent)
```

```
[1] 47.82975
```

median<mean - data appears to have a slight left skew.

Q10. Is the center of the distribution above or below 50%?

median<mean - center is below 50%.

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
#need to find winpercent value for chocolate candy?  
  
choc<- candy$chocolate==1 #defining chocolate candy  
fruit<-candy$fruit==1 #defining fruit candy  
  
choc.winp<-candy[choc,]$winpercent  
fruit.winp<-candy[fruit,]$winpercent  
  
choc.winp.avg<-mean(choc.winp) #mean of choc candy winpercent col.  
fruit.winp.avg<-mean(fruit.winp) #mean of fruit candy winpercent col.  
  
choc.winp.avg>fruit.winp.avg
```

```
[1] TRUE
```

Yes it is.

Q12. Is this difference statistically significant?

```
t.test(choc.winp,fruit.winp)
```

Welch Two Sample t-test

```
data: choc.winp and fruit.winp
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

The difference **is** statistically significant.

#3. Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

```
library(dplyr)

candy.desc<-candy%>%
  arrange(winpercent)%>%
  head(5)
candy.desc
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0	0	0	0
Boston Baked Beans	0	0	0	1	0	0
Chiclets	0	1	0	0	0	0
Super Bubble	0	1	0	0	0	0
Jawbusters	0	1	0	0	0	0
	crisped	rice	wafer	hard	bar	pluribus
Nik L Nip	0	0	0	1	0.197	0.976
Boston Baked Beans	0	0	0	1	0.313	0.511
Chiclets	0	0	0	1	0.046	0.325
Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511
	winpercent					
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					

The 5 least popular candy are Nik L Nip, Boston Baked Beans, Chiclets, Super Bubble, and Jawbusters

Q14. What are the top 5 all time favorite candy types out of this set?

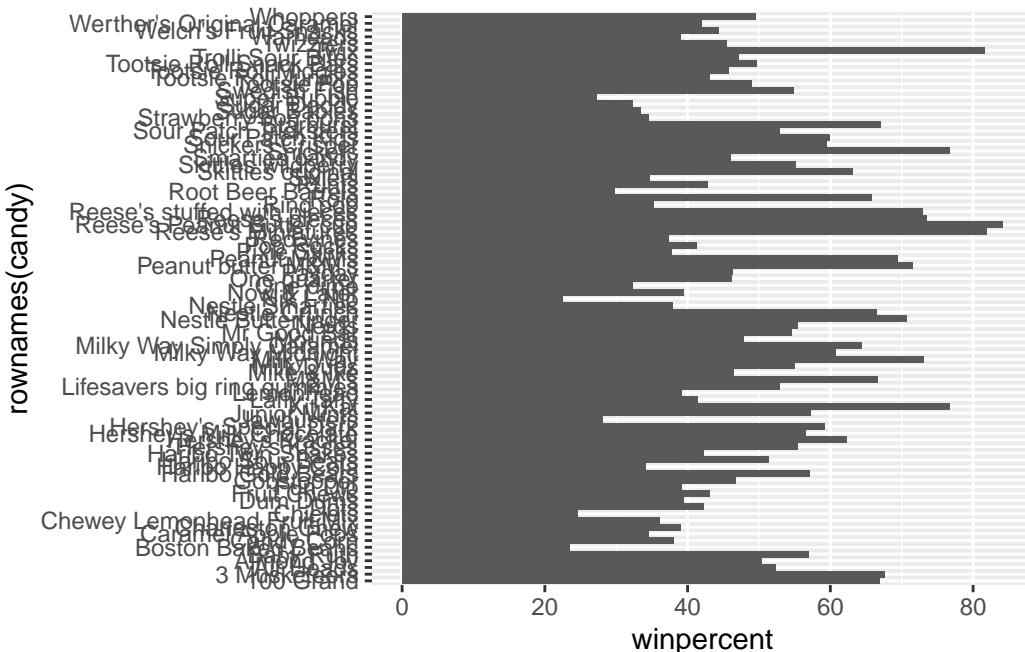
```
candy.asc<-candy%>%
  arrange(winpercent)%>%
  tail(5)
candy.asc
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Snickers	1	0	1		1	1
Kit Kat	1	0	0		0	0
Twix	1	0	1		0	0
Reese's Miniatures	1	0	0		1	0
Reese's Peanut Butter cup	1	0	0		1	0
	crispedrice	wafers	hard bar	pluribus	sugar	percent
Snickers	0	0	1	0	0	0.546
Kit Kat	1	0	1	0	0	0.313
Twix	1	0	1	0	0	0.546
Reese's Miniatures	0	0	0	0	0	0.034
Reese's Peanut Butter cup	0	0	0	0	0	0.720
	price	percent	winpercent			
Snickers	0.651	76.67378				
Kit Kat	0.511	76.76860				
Twix	0.906	81.64291				
Reese's Miniatures	0.279	81.86626				
Reese's Peanut Butter cup	0.651	84.18029				

The 5 most popular are Reese's Peanut Butter cups, Reese's Miniatures, Twix, Kit Kat, Snickers.

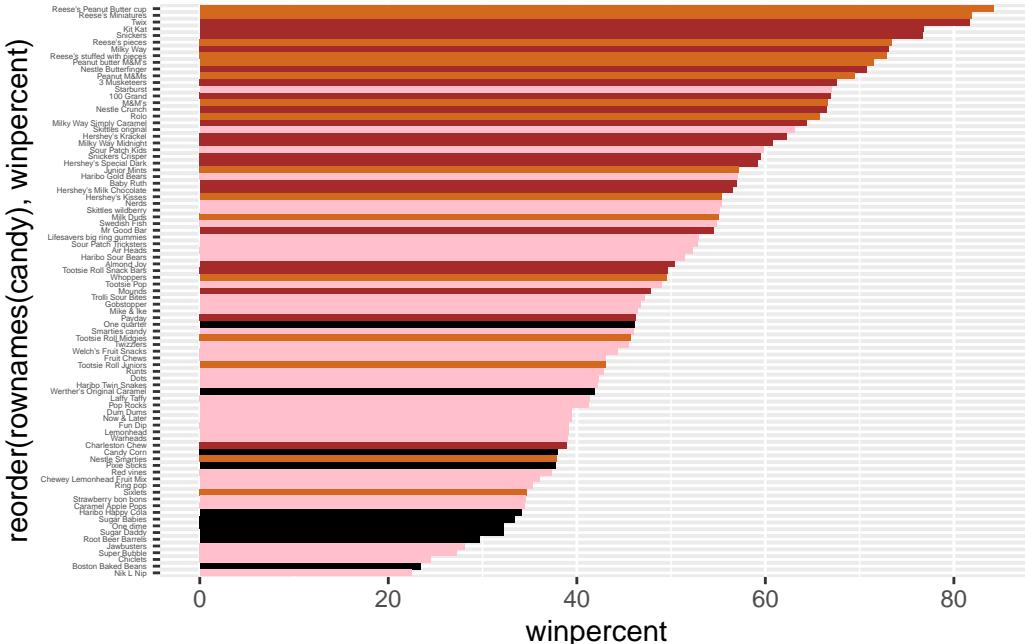
Q15. Make a first barplot of candy ranking based on winpercent values.

```
#simple plot not ordered
ggplot(candy,aes(winpercent,rownames(candy)))+
  geom_bar(stat="identity")
```



```
#defining colours
my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "pink"

#plot organised by colour
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=my_cols) +
  theme(axis.text.y = element_text(size = 3))
```



Now, for the first time, using this plot we can answer questions like: > Q17. What is the worst ranked chocolate candy?

Sixlets

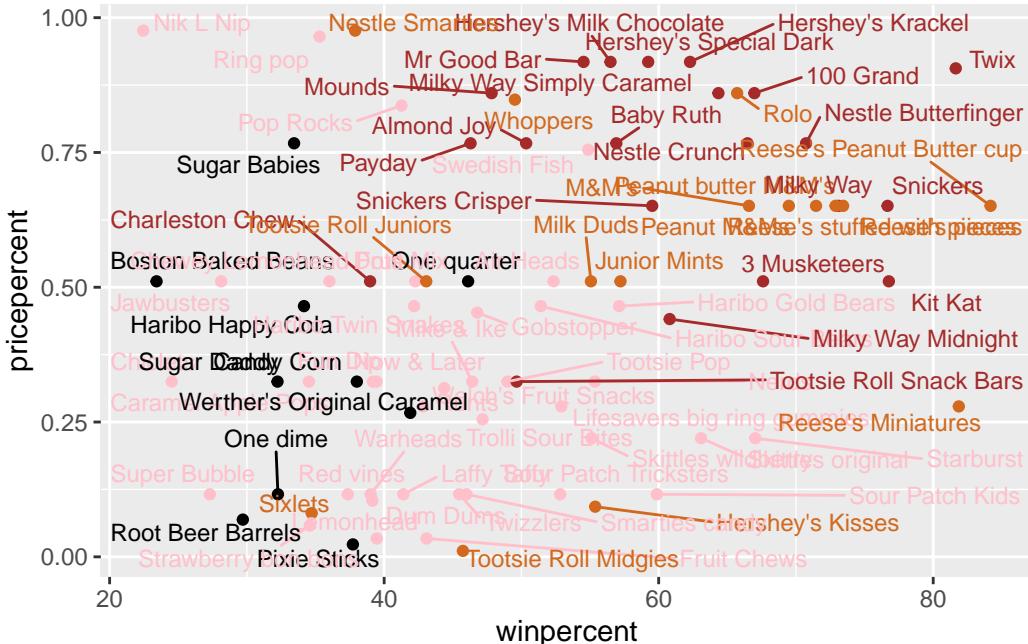
Q18. What is the best ranked fruity candy?

Starburst

4. Taking a look at pricepercent

```
library(ggrepel)

# How about a plot of price vs win
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 30)
```



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

Reese's Miniatures

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
ord <- order(candy$pricepercent, decreasing = TRUE)
head(candy[ord,c(11,12)], n=5)
```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Nik L Nip, Nestle Smarties, Ring pop, Hershey's Krackel, Hershey's Milk Chocolate

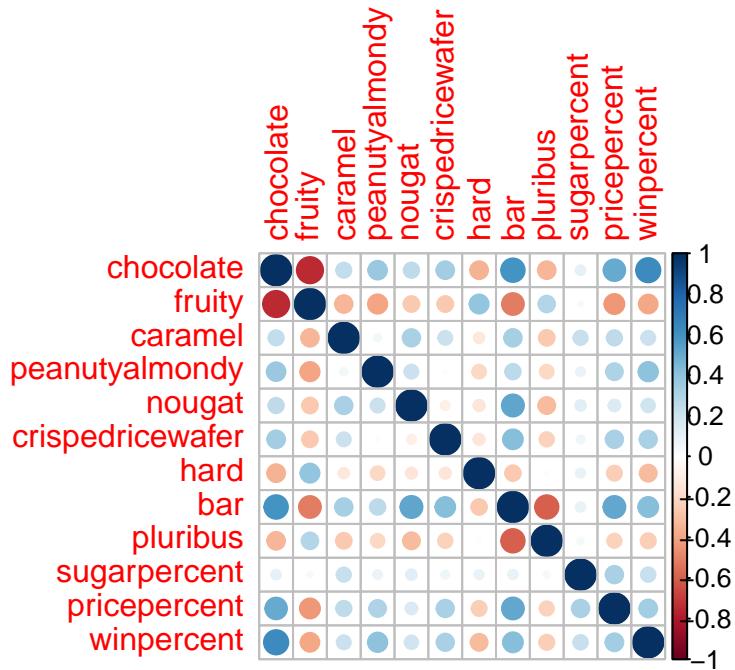
OPTIONAL Q21. Make a barplot again with geom_col() this time using pricepercent and then improve this step by step, first ordering the x-axis by value and finally making a so called “dot chart” or “lollipop” chart by swapping geom_col() for geom_point() + geom_segment().

5 Exploring the correlation structure

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
cij <- cor(candy)
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Fruity and Chocolate

Q23. Similarly, what two variables are most positively correlated?

Bar and Chocolate, Chocolate and Winpercent. I am choosing to not count the self comparisons (dark blue circles) since these will be 1 by default.

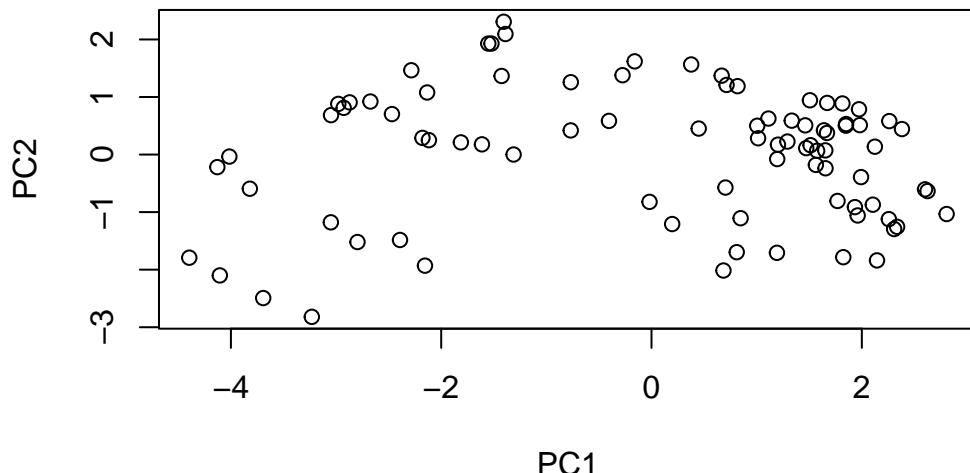
6. Principal Component Analysis

```
pca <- prcomp(candy, scale=TRUE)
summary(pca)
```

Importance of components:

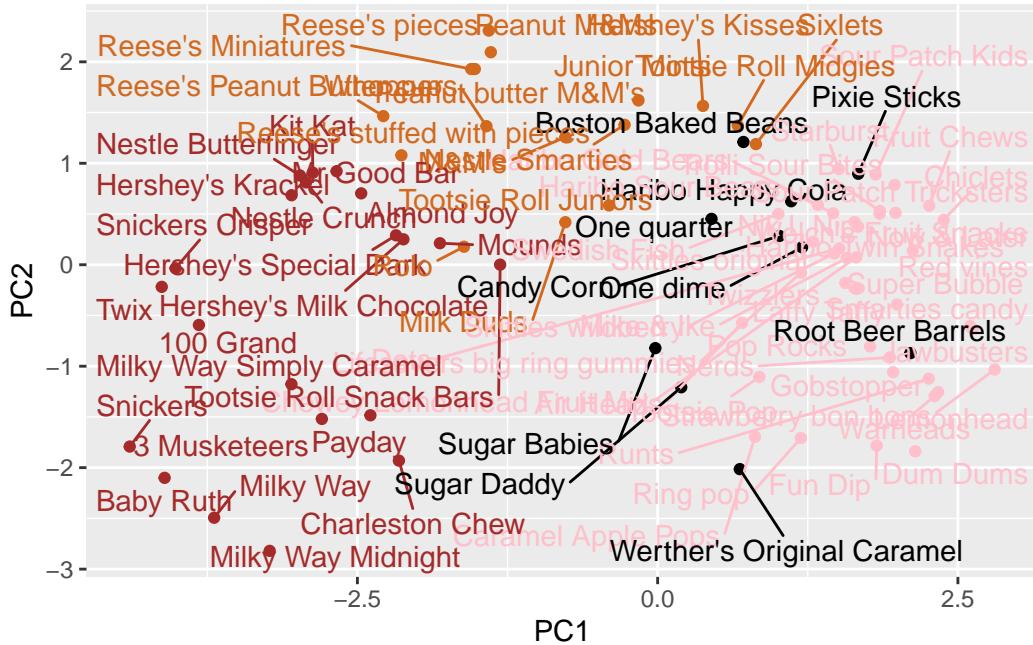
	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

```
plot(pca$x[,1:2])
```



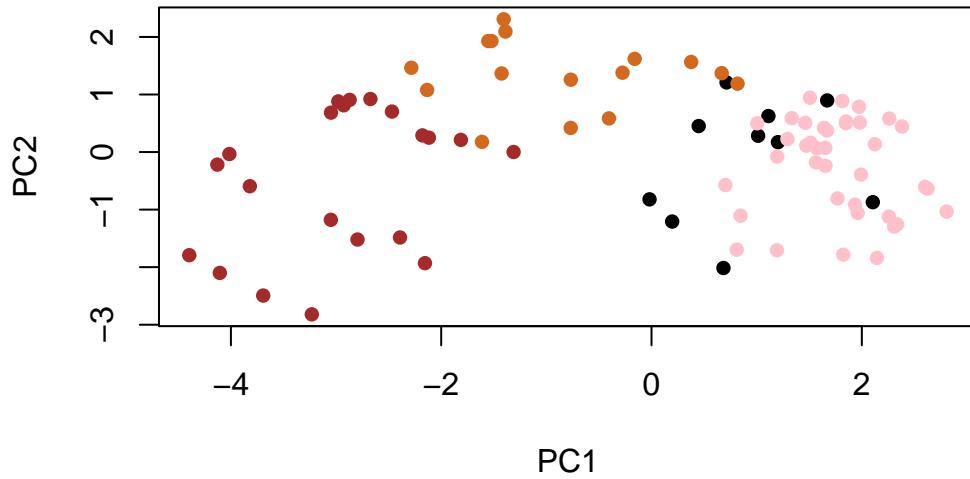
```
pc1<-pca$x[,1]
```

```
ggplot(pca$x, aes(PC1, PC2, label=rownames(pca$x)))+
  geom_point(col=my_cols)+
  geom_text_repel(col=my_cols, max.overlaps=100) #increased max overlap to see everything - b
```



We can change the plotting character and add some color:

```
plot(pca$x[,1:2], col=my_cols, pch=16)
```

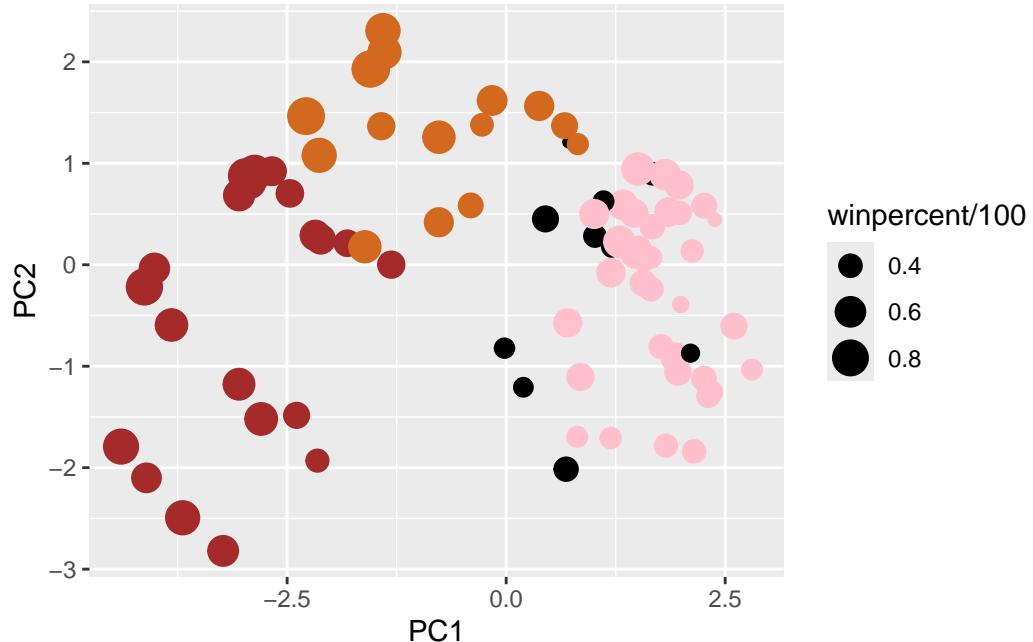


nicer plot with ggplot

```
# Make a new data-frame with our PCA results and candy data
my_data <- cbind(candy, pca$x[,1:3])

p <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=my_cols)

p
```

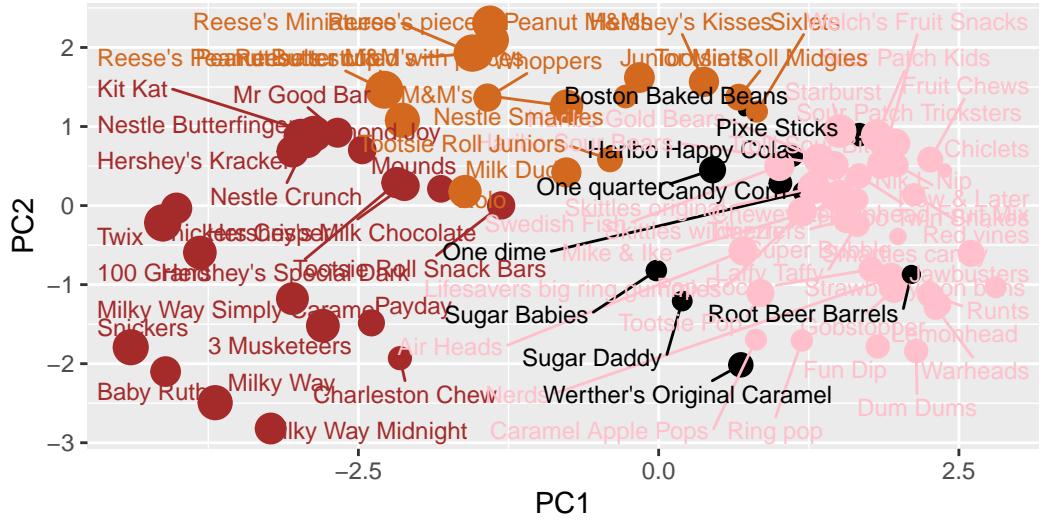


adding in labels

```
p + geom_text_repel(size=3.3, col=my_cols, max.overlaps = 50) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
       subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown),
       caption="Data from 538")
```

Halloween Candy PCA Space

Colored by type: chocolate bar (dark brown), chocolate other (light brown),



using `plotly` to generate an interactive plot

```
library(plotly)
```

```
Attaching package: 'plotly'
```

```
The following object is masked from 'package:ggplot2':
```

```
last_plot
```

```
The following object is masked from 'package:stats':
```

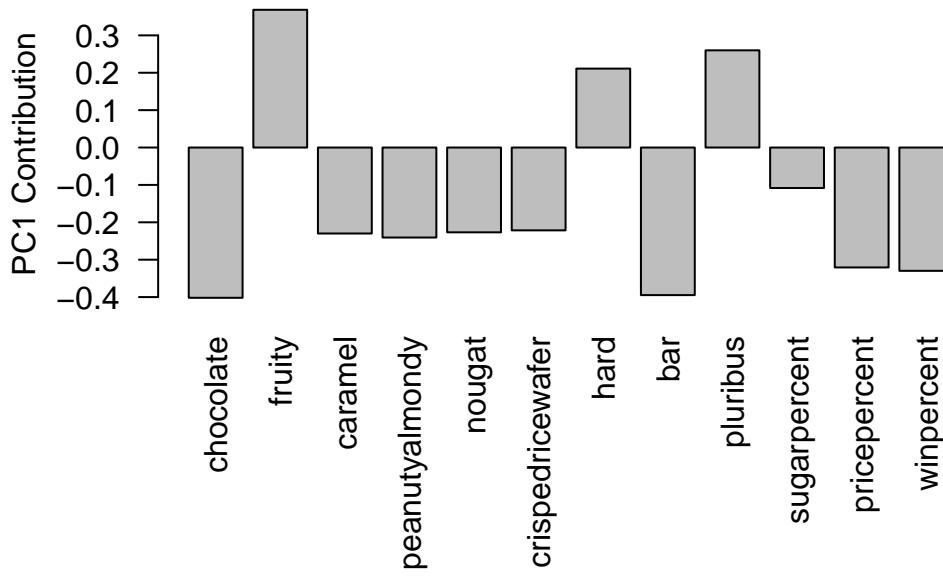
```
filter
```

```
The following object is masked from 'package:graphics':
```

```
layout
```

```
#ggplotly(p)

par(mar=c(8,4,2,2))
barplot(pca$rotation[,1], las=2, ylab="PC1 Contribution")
```

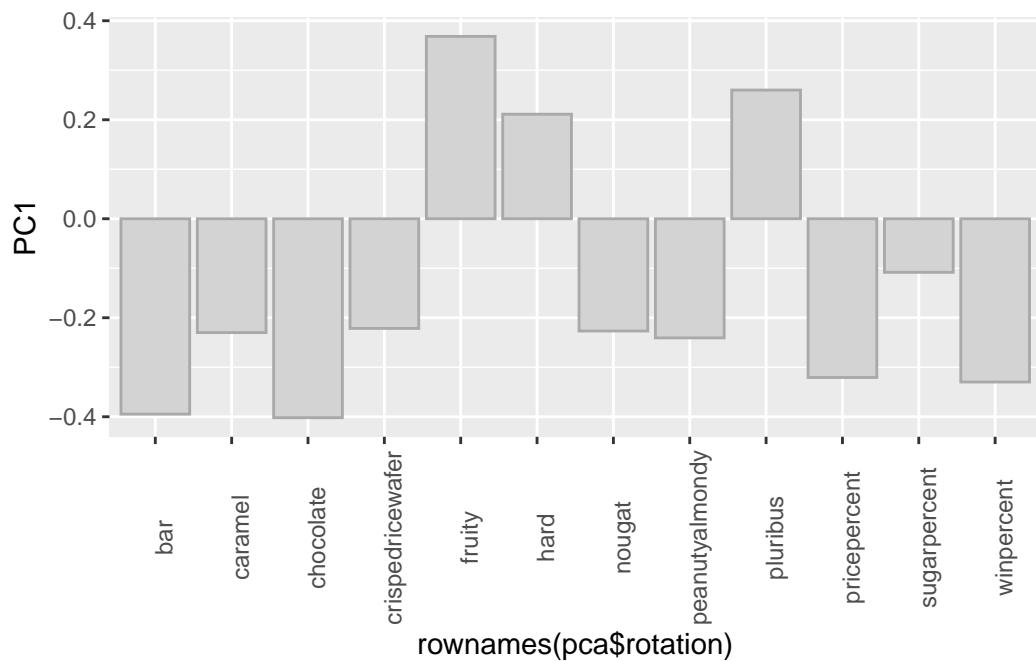


```
pca$rotation
```

	PC1	PC2	PC3	PC4	PC5
chocolate	-0.4019466	0.21404160	-0.01601358	0.016673032	-0.066035846
fruity	0.3683883	-0.18304666	0.13765612	0.004479829	-0.143535325
caramel	-0.2299709	-0.40349894	0.13294166	0.024889542	0.507301501
peanutyalmondy	-0.2407155	0.22446919	-0.18272802	-0.466784287	-0.399930245
nougat	-0.2268102	-0.47016599	-0.33970244	-0.299581403	0.188852418
crispedricewafer	-0.2215182	0.09719527	0.36485542	0.605594730	-0.034652316
hard	0.2111587	-0.43262603	0.20295368	0.032249660	-0.574557816
bar	-0.3947433	-0.22255618	-0.10696092	0.186914549	-0.077794806
pluribus	0.2600041	0.36920922	0.26813772	-0.287246604	0.392796479
sugarpercent	-0.1083088	-0.23647379	0.65509692	-0.433896248	-0.007469103
pricepercent	-0.3207361	0.05883628	0.33048843	-0.063557149	-0.043358887
winpercent	-0.3298035	0.21115347	0.13531766	-0.117930997	-0.168755073
	PC6	PC7	PC8	PC9	PC10
chocolate	0.09018950	0.08360642	-0.49084856	0.151651568	0.107661356
fruity	0.04266105	-0.46147889	0.39805802	0.001248306	0.362062502
caramel	0.40346502	0.44274741	0.26963447	-0.019186442	0.229799010
peanutyalmondy	0.09416259	0.25710489	0.45771445	-0.381068550	-0.145912362
nougat	-0.09012643	-0.36663902	-0.18793955	-0.385278987	0.011323453
crispedricewafer	0.09007640	-0.13077042	0.13567736	-0.511634999	-0.264810144

hard	0.12767365	0.31933477	-0.38881683	-0.258154433	0.220779142
bar	-0.25307332	-0.24192992	-0.02982691	-0.091872886	-0.003232321
pluribus	-0.03184932	-0.04066352	-0.28652547	-0.529954405	0.199303452
sugarpercent	-0.02737834	-0.14721840	-0.04114076	0.217685759	-0.488103337
pricepercent	-0.62908570	0.14308215	0.16722078	0.048991557	0.507716043
winpercent	0.56947283	-0.40260385	-0.02936405	0.124440117	0.358431235
	PC11	PC12			
chocolate	-0.10045278	-0.69784924			
fruity	-0.17494902	-0.50624242			
caramel	-0.13515820	-0.07548984			
peanutyalmondy	-0.11244275	-0.12972756			
nougat	0.38954473	-0.09223698			
crispedricewafer	0.22615618	-0.11727369			
hard	-0.01342330	0.10430092			
bar	-0.74956878	0.22010569			
pluribus	-0.27971527	0.06169246			
sugarpercent	-0.05373286	-0.04733985			
pricepercent	0.26396582	0.06698291			
winpercent	0.11251626	0.37693153			

```
ggplot(pca$rotation,aes(rownames(pca$rotation),y=PC1))+
  geom_bar(stat="identity",col="darkgrey",fill="lightgrey")+
  theme(axis.text.x = element_text(angle = 90))
```



Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

Fruity is the strongest positive variable in PC1. Other positive variables that are not as strong are hard and pluribus. This means that these variables correlate together, which makes logical sense since most fruity candy is hard and come in multiples. As PC1 value goes up (further along the x axis), these variables increase.